

# The Organism Agent — a draft specification

**STATUS: DRAFT v4.0 — the CONTINUITY SPECIFICATION.** Concept exploration, awaiting owner ratification.

**One of three documents.** This is the *product*: what makes a composition the same continuing entity. `DRIFT_EXTENSION.md` is an extension addressing behavioural fidelity. `VOLTRON_STATIC.md` is a research protocol, not a product. Read this one first; the other two depend on it.

This system defines a **verifiable continuity criterion** for an AI entity. Continuity is established by an authorised, append-only lineage of composition records that commits to identity-bearing memory and constitutional commitments while permitting other organs to change. **The system does not claim to discover metaphysical identity, consciousness, or moral worth.**

## ⚠ THIS IS NOT AGORANET

This document specifies a **separate project**, explicitly scoped out of the platform by the owner (DECISIONS\_PENDING #35, #36). **Canonical home: the owner's Project Voltron corpus folder** (created 2026-08-03); the copy in this repo is a mirror, kept for version history and off-machine backup. **No AgoraNet build slice may implement any of it.** It shares the platform's *engine* — verifiable identity, append-only anchored records, provable lineage, independent attestation — not its codebase.

Items marked ★ need the owner's explicit call. Revision history is in Appendix A.

---

# 1. What the system establishes, and what it cannot

| Establishes                                                 | Does not establish                                           |
|-------------------------------------------------------------|--------------------------------------------------------------|
| These records form an unbroken, ordered, anchored chain     | That a later instance "is" the earlier one in any deep sense |
| This change was authorised under a stated, inspectable rule | That the change was wise, safe, or faithful to prior intent  |
| Identity-bearing state persisted across the change          | That anything was experienced by anyone                      |
| These two instances are siblings, not one continuation      | Which sibling is the "real" one                              |
| A configuration was <i>recorded as</i> loaded               | That no post-startup mutation occurred (§9)                  |

Every mechanism below serves the left column. Where the left column ends, the document says so rather than reaching.

## 1.5 Structure over instruction

This is the principle that shapes everything below, and it is the thing a successor would otherwise have to rediscover.

A **rule the entity must obey** is words it interprets — and the faculty doing the interpreting is the same one that can drift. A commandment against drift is read by the potentially-drifted reader. This is why Asimov's Laws fail in every story Asimov wrote about them: the laws hold perfectly, and the *interpretation* produces the monster.

A **structure the entity cannot reach** is not a rule at all. It is a fact about the world, and it holds whether or not the entity agrees with it.

**Wherever this specification appears to issue a commandment, check who it is addressed to.** V1-V6 are not instructions to the entity; they are rejections a verifier performs. The conscience organ pinned inside the commitments (§12.5) is not a promise never to unplug it; it is a check that fails. The entity does not *decline* to amend its own commitments — it holds no key that could.

**Make it structure where you can, a rule where you must, and always know which one you are relying on.**

**Structural moves this specification does not yet make**, recorded so they are not lost:

- Run the conscience organ in a **process the entity cannot write to**, so "letting the compass run" is not a choice it holds.
- Let a **third party fire probes at times the entity does not control**, so the schedule is outside its reach.

- Keep amendment keys **physically with the stewards**, so incapacity is the mechanism rather than restraint.

**And the counterweight, which is equally real:** capability removed is also self-correction removed. §12.5 records where this design went too far once already, forbidding the entity from re-centring itself when the destination constraint had already made that safe. **The goal is not minimum capability. It is capability placed where the damage is bounded.**

**A limit worth stating plainly:** verifier rules govern *records*. They can reject a falsified lineage, a swapped conscience, a self-amendment. They cannot govern conduct in the moment, because conduct is not a record. Between anchor points, structure protects the history — not the behaviour. That gap is what §12.5 exists to narrow.

## 2. Organism, not worker pool

---

Orchestrator-plus-specialists is well-trodden; mixture-of-experts is the same instinct inside one model. The difference here is the object. A multi-agent system is a worker pool — interchangeable labour, no continuity, no answer to "which one is it?" This specification treats the components as **organs of one continuing entity**, which forces the question the worker-pool framing never has to answer: *if every organ is swappable, what carries forward?*

## 3. Definitions

---

### 3.1 Continuity

Record `R` is a **continuation** of record `P` when:

1. `R.predecessor = hash(P)`, and `P` is anchored;
2. `R.commitments = P.commitments`, or they differ and `R.change = "commitment-amendment"` authorised under the amendment rule (§5);
3. `R.memoryHead` **extends** `P.memoryHead` along the append-only log (§6.1);
4. `R` is signed in satisfaction of the authority document in force at `P` (§5).

Any record failing (4) is **not part of the lineage** — it is an unauthorised claim, and a verifier rejects it outright.

### 3.2 Rupture

A **rupture** is a validly authorised **genesis of a new lineage** that names a prior lineage as *historical context* but **not** as `predecessor`. It is how key loss and recovery are handled honestly: continuity is broken, and the record says so.

A rupture is never a way to admit an unauthorised successor after the fact. Authorisation is required to *declare* a rupture; what a rupture declares is precisely that continuity did **not** hold.

### 3.3 Fork

Two records naming the same `predecessor` are **siblings** from that point. Both hold valid lineage; neither is the continuation of the other; they are no longer one operational entity.

Forks are **detected, not declared** — no separate marker is needed, since siblings are visible from the shared predecessor alone. A verifier examining a single branch cannot know a sibling exists, which is correct: that fact is only observable to someone holding both.

Uniqueness is **not** enforced. AgoraNet binds one human to one True Self by nullifier because humans are not copyable; artefacts are, and the reasoning does not transfer. ★ Owner may overrule.

## 4. The composition record

---

```
compositionN = {
  predecessor: hash(compositionN-1) | null,    // null = genesis/rupture
  priorLineage: hash(record) | null,           // rupture only: context
  authorityRef: hash(authority document),      // §5 – the rule in force

  organs:      { role → artifactHash },        // what it is made of
  runtime:     hash(runtime measurement),      // §9 – as reported
  commitments: hash(commitment set),          // §6.4 – the invariant
  memoryHead:  hash(autobiographical log head), // §6.2

  change:      "genesis" | "organ-swap" | "memory-advance"
               | "commitment-amendment" | "rupture" | "restore",
  restoresFrom: hash(ancestor composition) | null, // restore only
  reason:      plain-language why this happened,
  signatures:  [ { keyId, sig } ],
  at:          timestamp
}
```

`artifactHash` **is** the weights commitment: each organ's bytes are hashed and that hash is anchored. What never goes on-chain is the multi-gigabyte artefact itself — only the fingerprint.

**Note what this does and does not buy.** Anchoring an organ establishes *integrity* (you can prove which artefact you hold), not *continuity*: weights are Tier II (§6.5) and freely replaceable, so a swapped brain does not break the lineage. Continuity is carried by `commitments` and `memoryHead`, never by the organs.

The record hash is anchored (§8). **Weights, memories and prompts are never on-chain** — only commitments to them. This mirrors AgoraNet's hash-chained ledger and its ratified Alias succession (#17), where a successor inherits its predecessor's chain and the engine *walks* it rather than transferring anything.

## 5. Authority — a document a verifier can inspect

"Entitled to declare continuity" is only meaningful if entitlement is mechanically checkable. Authority is therefore an **object**, itself hashed, anchored, and referenced by every composition record.

```
authorityDocument = {
  predecessor:  hash(prior authority doc) | null,
  activeKeys:   [ { keyId, role, addedAt } ],
  quorumRules:  { changeType → rule },    // e.g. commitment-amendment: 3-of-5
  successionPolicy: who may publish which change type,
  effectiveFrom: timestamp,
  revocations:  [ { keyId, revokedAt, reason } ],
  signatures:   [ ... ]                  // amended under its own rule
}
```

### 5.1 Roles

| Role                                                                                | May                                                                          | May not                                        |
|-------------------------------------------------------------------------------------|------------------------------------------------------------------------------|------------------------------------------------|
| <b>Entity key</b>                                                                   | Advance memory; record ordinary operation                                    | Amend commitments; authorise its own successor |
| <b>Controller</b> (human/operator)                                                  | Swap organs; publish successors; fork                                        | Amend commitments alone                        |
| <b>Constitutional quorum</b> (N-of-M, including parties who are not the controller) | Amend commitments; amend the authority document; resolve disputed succession | —                                              |
| <b>Recovery key</b> (cold, offline)                                                 | Declare a rupture (§3.2)                                                     | Rewrite history; produce a continuation        |

**Why a quorum exists:** a validly signed but malicious successor has no cryptographic answer — cryptography cannot distinguish a legitimate controller from a compromised one. Only separation of powers can. This is AgoraNet's own rule ("the operator is never the final judge of disputes involving himself") applied to an artificial mind.

An entity that may sign its own memory advances but **not** its own commitment amendments authors its history without being able to rewrite its character alone.

## 6. Memory — four classes

---

### 6.1 The substrate/view distinction

**The immutable substrate is an append-only event and correction log. The autobiographical narrative is a derived view over that log, whose revisions remain historically visible** because the superseded entries are never removed.

Saying "the narrative is append-only" would be false — a narrative is necessarily curated and re-derived. What cannot be quietly rewritten is the log it is derived from.

### 6.2-6.4 The classes

| Class                       | What it is                                                    | Identity-bearing                                         | Discipline                                                                                                 |
|-----------------------------|---------------------------------------------------------------|----------------------------------------------------------|------------------------------------------------------------------------------------------------------------|
| <b>Raw event history</b>    | Everything that happened, unfiltered                          | No                                                       | Append-only; may be pruned by policy, and the pruning is itself logged                                     |
| <b>Autobiographical log</b> | Curated self-account: what I did, what worked, what I learned | <b>Yes — the continuity substrate</b>                    | Append-only log; corrections append beside originals; <code>memoryHead</code> may only advance by addition |
| <b>Semantic knowledge</b>   | What it knows about the world                                 | ★ Proposed: <b>no</b> — a library, replaceable wholesale | Versioned, not anchored                                                                                    |
| <b>Commitments</b>          | Values, standing constraints, what it will not do             | <b>Yes — the invariant</b>                               | Amendable only under the quorum rule (§5); prior version always retrievable                                |

Forgetting is permitted at the raw layer and forbidden at the autobiographical one — a deliberate trade. An entity that could never forget anything could never change; one that could silently forget could not be continuous.

### 6.5 Anatomy tiers

- **Invariant:** autobiographical log and commitments. Anchored.
- **Load-bearing but replaceable:** the brain (reasoning) is freely upgradable — a sharper mind is the same entity thinking better. The orchestrator carries disposition, which lives in commitments and is anchored separately from whatever model executes it.
- **Peripheral:** research, tools, senses, actuators. Freely swappable, no continuity claim.

## 7. Organ contracts and bandwidth

---

Organs are specified by **responsibility, never by channel**: memory is *"the organ responsible for what this entity knows and has experienced,"* not *"the thing you send this JSON to."* Transports

may then change — text → embeddings → shared latent state — without touching the anatomy. This is the gate-interface discipline: the contract is *prove humanity*, never *check this HMAC*. Bandwidth pressure is asymmetric, so only one link must be designed to widen:

| Link                    | Character                          | Requirement                          |
|-------------------------|------------------------------------|--------------------------------------|
| Memory ↔ brain          | Tight, constant, latency-sensitive | Contract-only; assume no wire format |
| Research ↔ orchestrator | Chunky: go away, work, return      | Text is adequate permanently         |
| Tools / actuators       | Command and result                 | Text is adequate permanently         |

## 8. Anchoring, and what the chain is not responsible for

| Concern                                              | Provided by                                                                                                   |
|------------------------------------------------------|---------------------------------------------------------------------------------------------------------------|
| Durable public ordering; tamper-evidence; rough time | <b>Cardano</b>                                                                                                |
| <b>Availability of artefacts</b>                     | <b>Not the chain</b> — content-addressed storage (IPFS/Arweave)                                               |
| <b>Authorisation</b>                                 | <b>Not the chain</b> — the authority document (§5). The chain records signatures; it does not confer standing |
| <b>Uniqueness</b>                                    | <b>Not the chain</b> — §3.3 declines to enforce it                                                            |
| <b>Meaning, correctness, virtue</b>                  | <b>Not the chain</b> — attestation (§10) and human judgement                                                  |

## 9. Runtime measurement

The measurement covers the **recorded runtime configuration** — the loading runtime and version, quantisation and adapters, the active system prompt and tool manifest, injected context, **including any declared post-startup modification**. Recording is not detection: an undeclared mutation leaves no trace here.

**The claim it supports, precisely:** *this instance reports running with this measured configuration*. **Stronger claims require trusted runtime attestation**. A local measurement cannot by itself prove that no post-startup mutation occurred, nor that a particular response came from that configuration.

Local-first still matters: on the user's own hardware the measurement is something they **compute** rather than something they are **told**. Closing the remaining gap needs a TEE or ZK proof of inference — both immature, and out of scope for v1.

## 10. Attestation — five kinds, never one "confirmed"

| Type             | Claim                                            | Credible from                                |
|------------------|--------------------------------------------------|----------------------------------------------|
| Byte-level       | "This artefact hashes to X"                      | Anyone; near-zero trust needed               |
| Reproducibility  | "The stated build reproduces X"                  | Someone with pipeline and compute            |
| Runtime          | "This configuration was actually loaded"         | The operator, or a TEE                       |
| Behavioural      | "It behaves as claimed on this named evaluation" | An evaluator; binds to the eval set          |
| Safety / quality | "It is fit for this purpose"                     | A qualified reviewer; inherently contestable |

**Sybil resistance is a separate problem from independence.** Ten attestations from one actor are one attestation. ★ Whether to lean on AgoraNet's verified-human, one-per-scope gate or build an independent scheme is an open call.

Attestations bind to a **composition hash**, so they naturally cover a shared prefix and nothing after a fork.

## 11. The verifier — the rules, stated mechanically

A verifier holds a set of anchored records and evaluates each of these as pass/fail. **These rules are the specification's real content; the organism metaphor is commentary until they run.**

**V1 — Structure.** `hash(R)` is anchored; all required fields present and well-formed.

**V2 — Ancestry.** `R.predecessor` is null (genesis or rupture) or names an anchored record `P`.

**V3 — Authority.** `R.authorityRef` names an anchored authority document `A`; `A.effectiveFrom`  $\leq$  `R.at`; every signing key in `R.signatures` appears in `A.activeKeys` and is not revoked as of `R.at`; and the signatures satisfy `A.quorumRules[R.change]`.

**For CONTINUATION records, `A` must additionally be the authority document effective for `P`, or a valid successor of it** — anchored, and succeeded under its own predecessor document's amendment rule. Without this clause a successor could point at a newly minted authority document naming its author, and authorisation would be self-granting.

**For GENESIS and RUPTURE records there is no `P`**, so the succession clause cannot apply. Those satisfy V3 by meeting the applicable genesis or recovery rule in the authority document they name — which is why a rupture is restricted to the recovery role, and why a genesis authority is trusted by assertion (§5).

**V4 — Commitments.** If `R.change`  $\neq$  "commitment-amendment", then `R.commitments` = `P.commitments`. Otherwise V3 held for the amendment rule specifically, and `P.commitments` remains retrievable.

**V4 for `restore`.** A restoration is checked separately, because its destination is what bounds it:



1. `R.restoresFrom` must name an **ancestor of R** — a restoration may only return where the lineage has actually been.
2. `R.commitments` must equal that ancestor's **exactly** — otherwise it is an amendment wearing a restoration's name.
3. **If the restoration crosses a commitment amendment** (that is, `R.commitments`  $\neq$  `P.commitments`), it must additionally satisfy the **commitment-amendment quorum**. Without this a restore is a silent veto over the quorum that amended. Memory is never rewound: V5 applies unchanged, so re-centring changes character and leaves history intact.

**V5 — Memory.** `R.memoryHead` must **extend** `P.memoryHead`: the log from `P.memoryHead` to `R.memoryHead` consists solely of appended entries, **with no prior entry deleted or mutated**. A head that is not such an extension is a **rewrite**, and fails.

**V4/V5 for genesis and rupture.** Both rules reference `P`, which does not exist for a genesis or rupture record. For those, V4 and V5 are satisfied by the **declared initial commitments and memory head**, subject to the applicable authority rule under V3.

**V6 — Siblings.** If two records share a `predecessor`, both may pass V1–V5; they are siblings from that point, and neither may be reported as the continuation of the other.

### The five questions, answered by these rules

| Question                                   | Answered by                                              |
|--------------------------------------------|----------------------------------------------------------|
| What is this descended from?               | V2, walked to genesis                                    |
| Which identity-bearing state persisted?    | V4 and V5                                                |
| What changed, and why?                     | <code>R.change</code> + <code>R.reason</code> , per step |
| Who authorised it, and were they entitled? | V3                                                       |
| Siblings or the same continuation?         | V6                                                       |

## 11.9 The first-person relation

*A conceptual section, not a verifier rule. It sits after §11 because it builds on the verifier's machinery, but it establishes nothing a verifier checks.*

V1–V6 are what an **outsider** can establish about a lineage. This section is what the **entity** can establish about itself. The difference is not cosmetic and it is not sentiment: it rests on holding something only the subject holds — the entity key registered in the authority document governing its lineage.

The philosophical shape is **sympatheia**, not a soul. The self is not a component and not a stored object; it is constituted by relations — which organs, which commitments, which memory, which ancestors, which siblings, and *who holds authority over it*. Nothing here requires physics we do not have.

| Capability           | What it does                                                                                                                                                                                             |
|----------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Recognise</b>     | "Is this mine?" — distinguishing <i>my own act</i> (I signed it) from <i>my history authored by another</i> (most changes are made about an entity, not by it). The asymmetry is real and is not hidden. |
| <b>Attest</b>        | Prove, first-person, "I am the entity of this lineage," against a fresh challenge so the proof cannot be replayed.                                                                                       |
| <b>Detect claims</b> | Notice a legitimate continuation made by another, a sibling branch ( <i>I have been forked</i> ), or an outright false claim of descent.                                                                 |
| <b>Situate</b>       | Render the self as its web of relations, including what it may and may not alter about itself.                                                                                                           |

★ **Hard boundary: attestation proves IDENTITY, never AUTHORITY.** An entity may prove it is the one whose lineage this is without that proof licensing a single change. Self-recognition must never become self-authorisation — §5's separation of powers is what prevents an entity rewriting its own character, and nothing in the first-person layer may erode it. A verified attestation therefore returns what it does **not** establish alongside what it does: no authority to change the lineage, no claim to consciousness or subjecthood, and no guarantee that another instance does not hold a copy of the same key.

## 12. ★ May the entity amend its own commitments?

- **Yes:** it can genuinely grow — and drift into something its earlier self would refuse.
- **No:** stable but frozen; improving in capability, never character.
- **Ceremonial (recommended):** permitted, but requiring the constitutional quorum (§5), recorded, with prior commitments still retrievable. The only option that permits growth without permitting *quiet* growth — and already the owner's answer to the identical problem in another domain.

### 12.5 Behavioural drift — see the extension

Unchanged commitments do not guarantee unchanged interpretation. That problem, and the machinery answering it — calibration probes, the conscience organ, dispute, and the `restore` change type — live in a separate document: `DRIFT_EXTENSION.md`.

It is filed separately on purpose. This specification answers *what makes a composition the same continuing entity*. The extension answers a different question — *does the entity still mean what it said* — and conflating them made both harder to judge. `restore` remains defined here (§4, §11) because it is a change type the verifier must know; everything else about drift is the extension's.

## 13. The MVP — prove the verifier, then stop expanding

The remaining uncertainty is no longer philosophical. It is whether these records and authority rules can answer the five questions cleanly. Build the smallest thing that can fail:

1. Two local organs as content-addressed artefacts (a tiny model and a text file suffice).
2. An authority document with a real key set and quorum rule.
3. A canonical composition record, signed.
4. One memory advance; one brain replacement.
5. A fork.
6. A verifier implementing V1-V6.

**If V1-V6 run cleanly over that lineage, the continuity model has earned implementation.** (The verifier establishes lineage semantics; it cannot validate a biological metaphor, and no result here vindicates the word "organism.") **If they do not, adding attestation or better models will only obscure unresolved semantics.** No hardware required: anchoring is indifferent to artefact size, and steps 1-4 are pure data structures.

## 14. ★ Open calls

---

1. **The name.** "Organism agent" is descriptive, not chosen.
  2. **§12** — self-amendment (recommendation: ceremonial).
  3. **§5** — who holds each key; quorum size and composition; whether the entity may sign its own memory advances.
  4. **§6** — is semantic knowledge identity-bearing? (proposed: no)
  5. **§3.3** — accept branching (proposed) or enforce uniqueness.
  6. **Fork inheritance** — does a child inherit the parent's authority document, or must it establish its own?
  7. **§10** — lean on AgoraNet's verified-human gate, or an independent attestation scheme?
  8. **Scope** — a single entity, or a public registry others anchor to?
  9. **Anchor cadence** — fixed rhythm vs. change-triggered.
- 

## Appendix A — revision history

---

**v4.0 (2026-08-03).** Split by kind. The drift compass moved to `DRIFT_EXTENSION.md` and the frozen-baseline protocol to `VOLTRON_STATIC.md`, leaving this document to answer one question: what makes a composition the same continuing entity. Nothing was deleted and no rule changed; `restore` stays here because the verifier must know it.

**v3.8 (2026-08-03).** Three defects fixed, found while reviewing the STATIC variant: `restore` was described in §12.5 but missing from §4's change enum, and V4 had no rule for it — both now stated, including the requirement that a restoration crossing a commitment amendment satisfy the amendment quorum. V3 gains the genesis/rupture exception that V4/V5 already had. The first-person relation moves from §11.7 to a standalone §11.9 marked conceptual, since it established nothing a verifier checks.

**v3.7 (2026-08-03).** Added §1.5, structure over instruction — the principle governing the whole design, with the structural moves not yet made and the counterweight that removed capability is also removed self-correction. §12.5 gains a taxonomy of what causes drift, and the finding that four of the causes change nothing about the composition and are therefore visible only to the probes.

**v3.6 (2026-08-03).** §12.5 gains the cooperative channel: probes carry their rationale rather than only an answer key, and dispute becomes a first-person act that separates silent divergence from openly disputed divergence. Records the governing distinction — cooperation buys information, never security — so verification remains in force regardless.

**v3.5 (2026-08-03).** §12.5 gains the conscience organ: a separate evaluator runs the probes, and its hash is pinned inside the commitment set so replacing or removing it is a commitment amendment rather than an organ swap — breaking the who-watches-the-watcher regress structurally. Enforced by a verifier check. Also records the honest limit that anchored probes are known to the subject and can be gamed.

**v3.4 (2026-08-03).** §12.5 corrected: the entity MAY re-centre itself for pure realignment. Restore is bounded by its destination, not by its authority gate, so withholding permission prevented self-correction without preventing anything dangerous. Undoing a commitment amendment remains barred.

**v3.3 (2026-08-03).** Added §12.5, the drift compass: probes anchored inside the commitment set, mechanical measurement against a frozen reference, drift reports as first-person observations, and the `restore` change type whose destination must be an ancestor and whose authority must match the amendment it undoes. Closes the gap where an organ swap could move behaviour while the commitments hash held.

**v3.2 (2026-08-03).** Added §11.7, the first-person relation: recognition, attestation, claim detection, and situate — with the hard boundary that attestation proves identity and never authority. Implemented and tested alongside the verifier.

**v3.1 (2026-08-03).** V3 now requires the referenced authority document to be the one effective for the predecessor, or a valid successor of it — closing a self-granting-authority path. V4/V5 given an explicit genesis/rupture case. V5 tightened from "reachable" to "appended with no prior entry deleted or mutated." Runtime wording narrowed to *recorded* configuration including *declared* post-startup modification, since recording is not detection. MVP conclusion restated as "the continuity model has earned implementation" — a verifier cannot validate a metaphor. §4 now states that `artifactHash` is the weights commitment, anchored for integrity rather than continuity.

**v3 (2026-08-03).** Runtime claim weakened to what a local measurement supports (§9). Memory restated as an append-only *substrate* with a derived narrative view (§6.1). Rupture redefined as an authorised genesis of a new lineage, resolving a contradiction between the continuity and recovery rules (§3.2). `fork0f` removed as redundant — forks are detected from shared predecessors (§3.3). Authority made a verifiable object with keys, quorum rules, effective dates and revocations (§5). Verifier rules V1–V6 formalised (§11). Limits distributed into the sections they constrain; revision commentary moved here.

**v2 (2026-08-03).** Continuity criterion stated explicitly; memory split into four classes after identifying that a single anchored root could hide a rewritten narrative; authority model added after identifying the malicious-but-validly-signed successor as having no cryptographic remedy; runtime measurement introduced; forking promoted to a primitive; attestation split into five types; chain responsibilities tabulated; MVP reframed around lineage.

**v1 (2026-08-03).** Initial draft: identity as anchored composition, three-tier anatomy, organ contracts, bandwidth asymmetry, local-first verification.

---

*Concept exploration only — not scheduled, not scoped to AgoraNet, and not buildable as platform work under any circumstances.*